

Access Code: **340362**

**CMP9794M**

# **Advanced Artificial Intelligence**

[Heriberto Cuayahuitl](#)



UNIVERSITY OF  
**LINCOLN**

School of Engineering and Physical Sciences



# Last Week

- Discussion on structure learning
- Algorithms to induce the structure of Bayes nets
  - PC-Stable (step 1 in particular)
  - Hill Climbing
  - Min-Max Hill Climbing
- Multiple metrics to measure the performance of probabilistic models on test data

# Today

- **Inference by stochastic simulation**
- Approximate inference in Bayesian nets (BNs)
  - Rejection Sampling
  - Likelihood Weighting
  - Gibbs Sampling
- Data discretisation for learning discrete BNs
- Questions and answers related to assignment

# Inference by Stochastic Simulation

- Basic idea:
  - Draw  $N$  samples from a sampling distribution  $S$ 
  - Compute an approximate posterior probability  $\hat{P}$ 
  - $\hat{P}$  converges to the true probability  $P$  as the number of samples tends to infinity.
- That can be used when exact inference is intractable (in the case of large Bayes nets or many hidden vars.)
- Let's look at an example of the basic idea.

# Inference by Stochastic Simulation

- Imagine we have a stochastic process that generates numbers in the range: 0, 1, . . . 9
- How do we determine the probability of each value being generated?
  - One possibility is to generate lots of numbers (samples) and look at the **sample average** for each value.
  - If we **generate  $N$  samples**, and  **$m$  of them are 9**, then the sample average is:  $\frac{m}{N}$
  - This is an estimate of the probability of 9 being generated.

# Inference by Stochastic Simulation

- Overall process:
  - Draw  $N$  samples from a sampling distribution  $S$ . The “sampling distribution” is generated by the process.
  - Compute an approximate posterior probability  $\hat{P}$ . It is an estimate (what the hat means) because it is approximate.
  - $\hat{P}$  converges to the true probability  $P$  as the number of samples tends to infinity.
- But note that **the more samples, the more accurate** our estimates will be.

# Inference by Stochastic Simulation

- This works for any process where we can just observe the samples:
  - Rain in Lincoln (in October).
  - How many students turn up to Advanced AI.
  - How many apples my apple tree produces.
  - How long it takes me to walk to work.
- However, for our BN we have to generate the samples.
- That is what we will look at next.

# Probabilistic Sampling: Example

- Consider random variable  $C$  with domain values {red, green, blue} and probability distribution:  $P(C = \text{red}) = 0.6$ ,  $P(C = \text{green}) = 0.1$ ,  $P(C = \text{blue}) = 0.3$
- Generating 8 random numbers looks like:



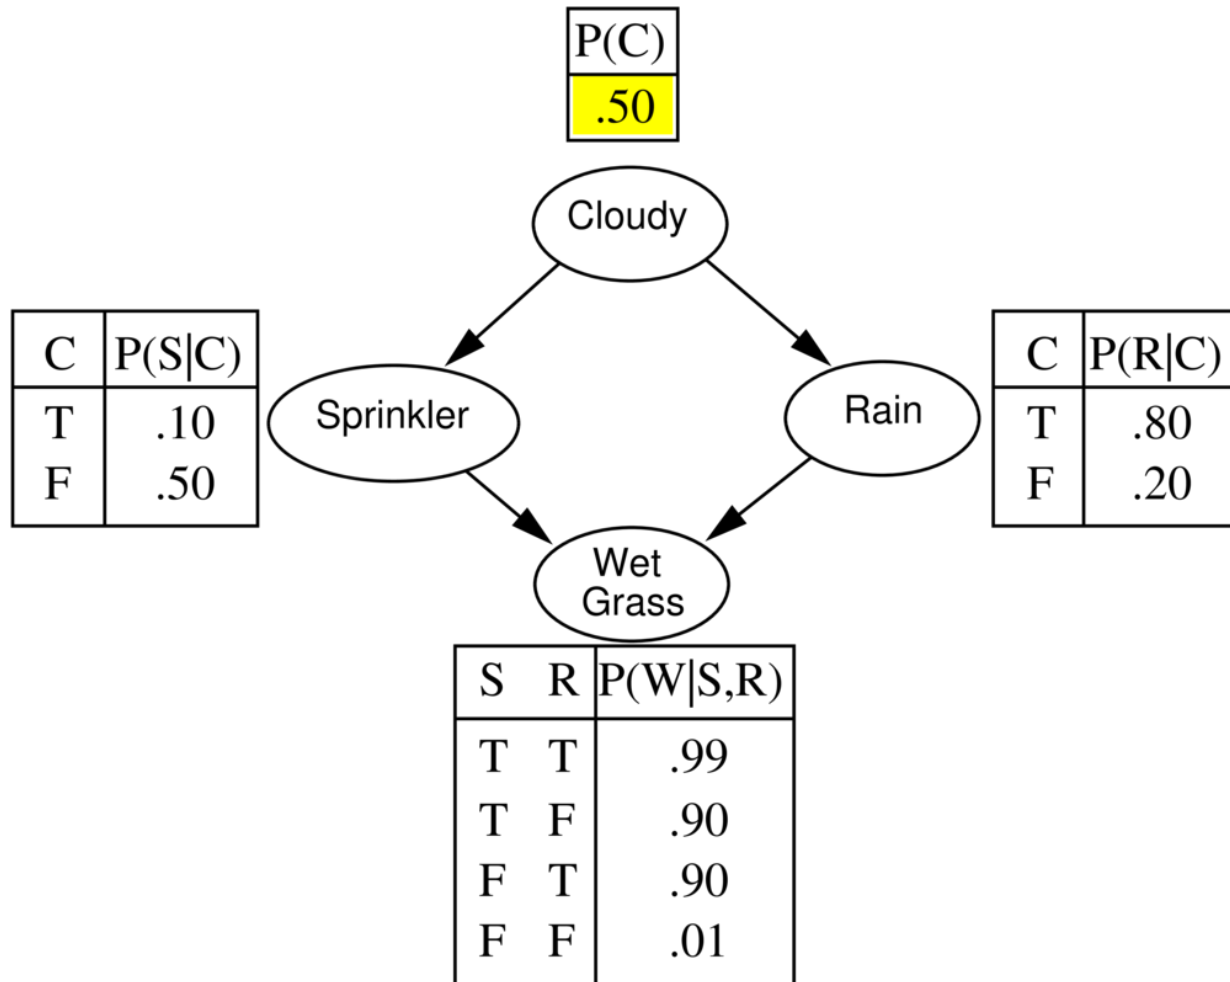
- That gives the following:  $P(C) = \langle \frac{5}{8}, \frac{1}{8}, \frac{2}{8} \rangle$



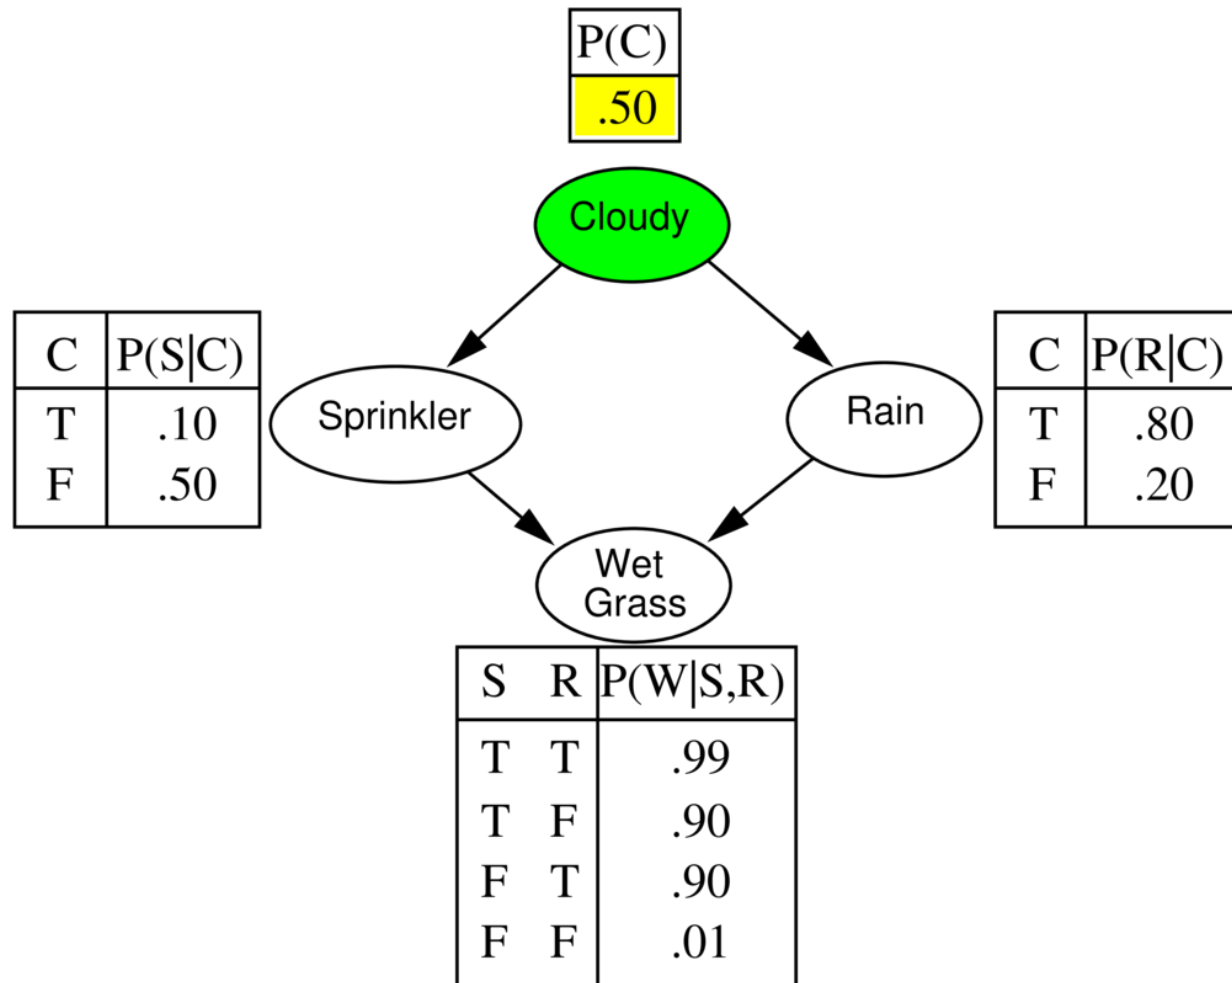
# Today

- Inference by stochastic simulation
- Approximate inference in Bayesian nets (BNs)
  - **Rejection Sampling**
  - Likelihood Weighting
  - Gibbs Sampling
- Data discretisation for learning discrete BNs
- Questions and answers related to assignment

# Prior Sampling

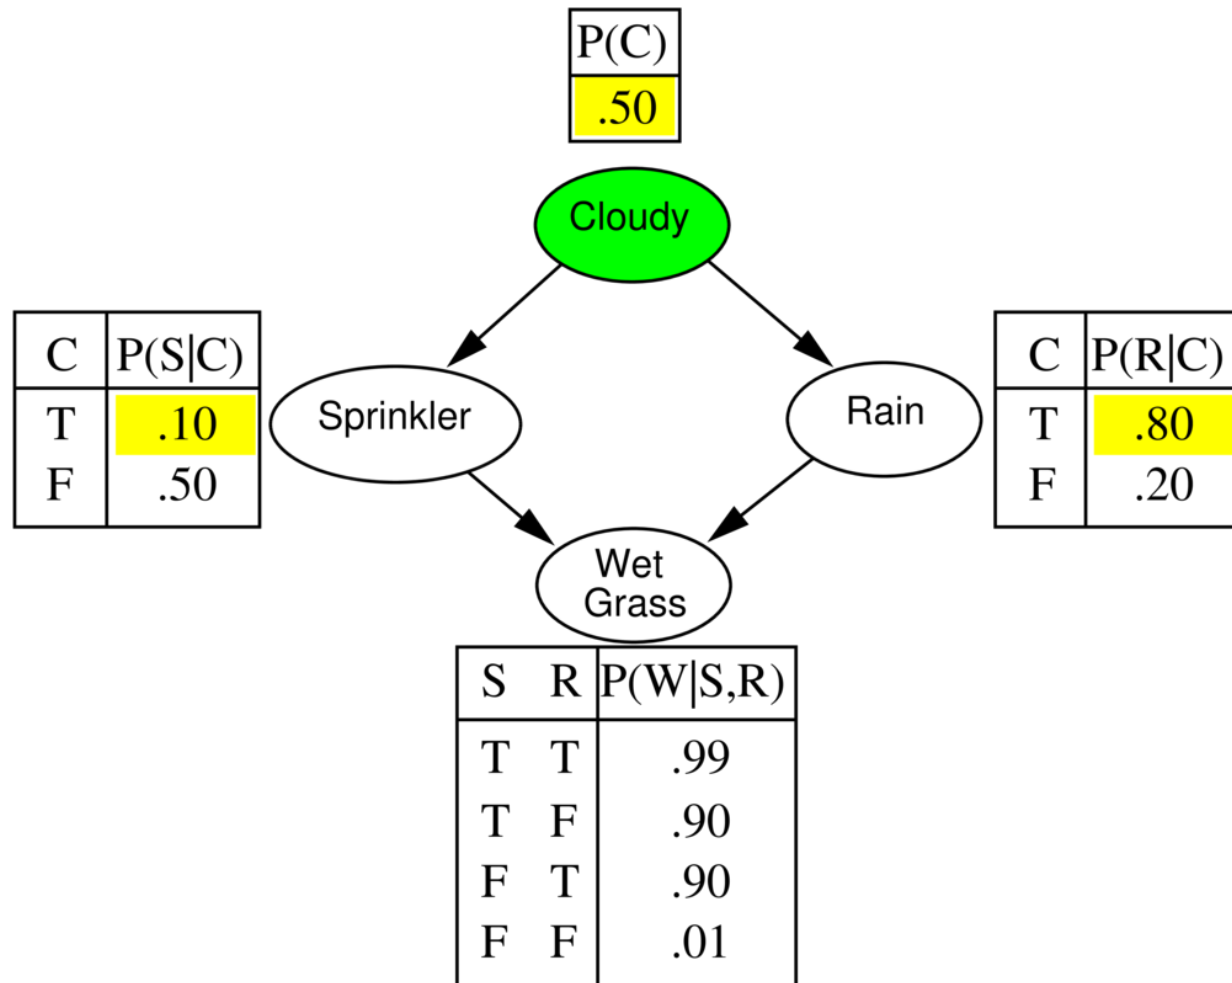


# Prior Sampling

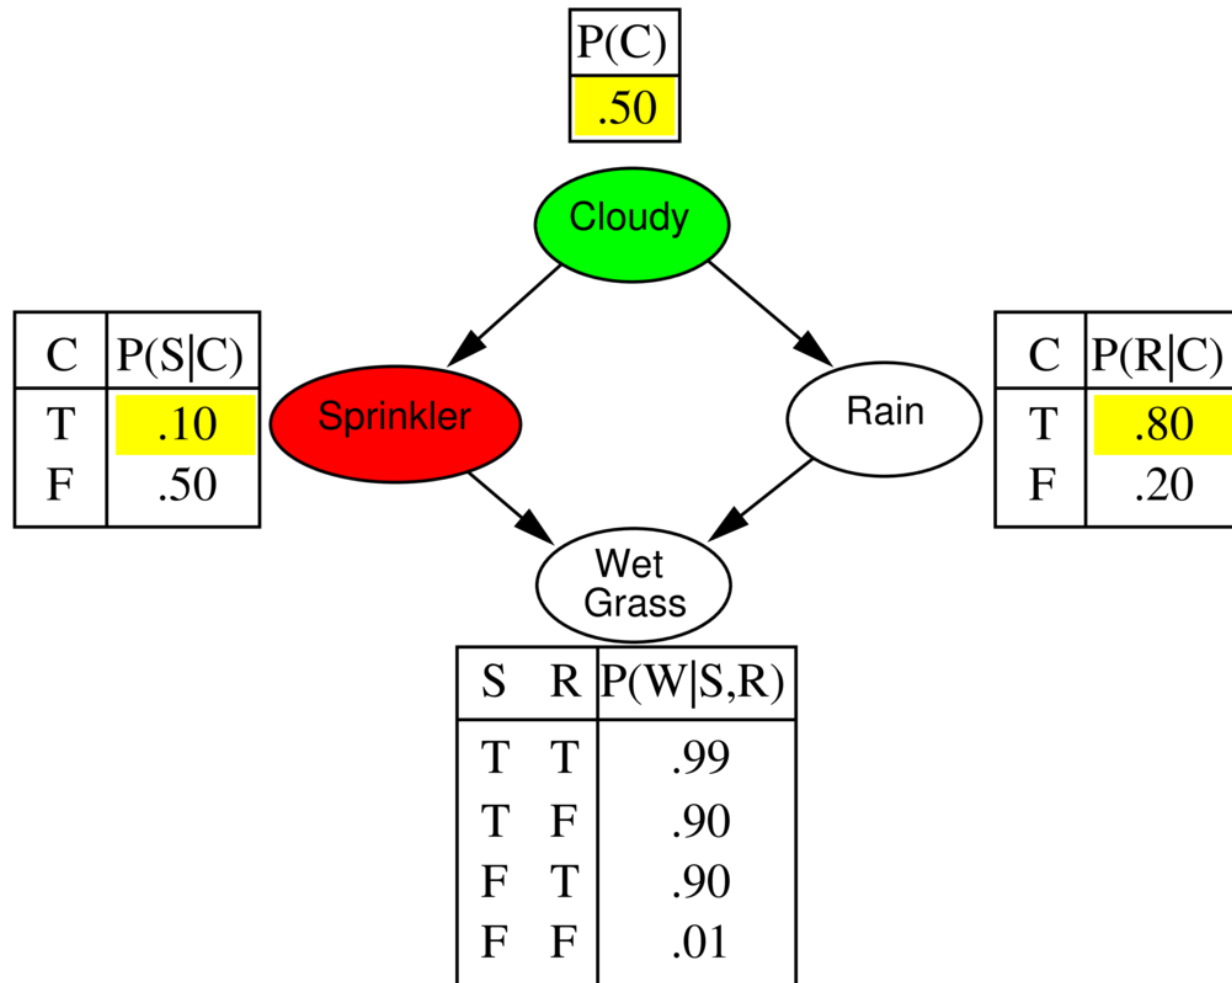


Cloudy = true

# Prior Sampling

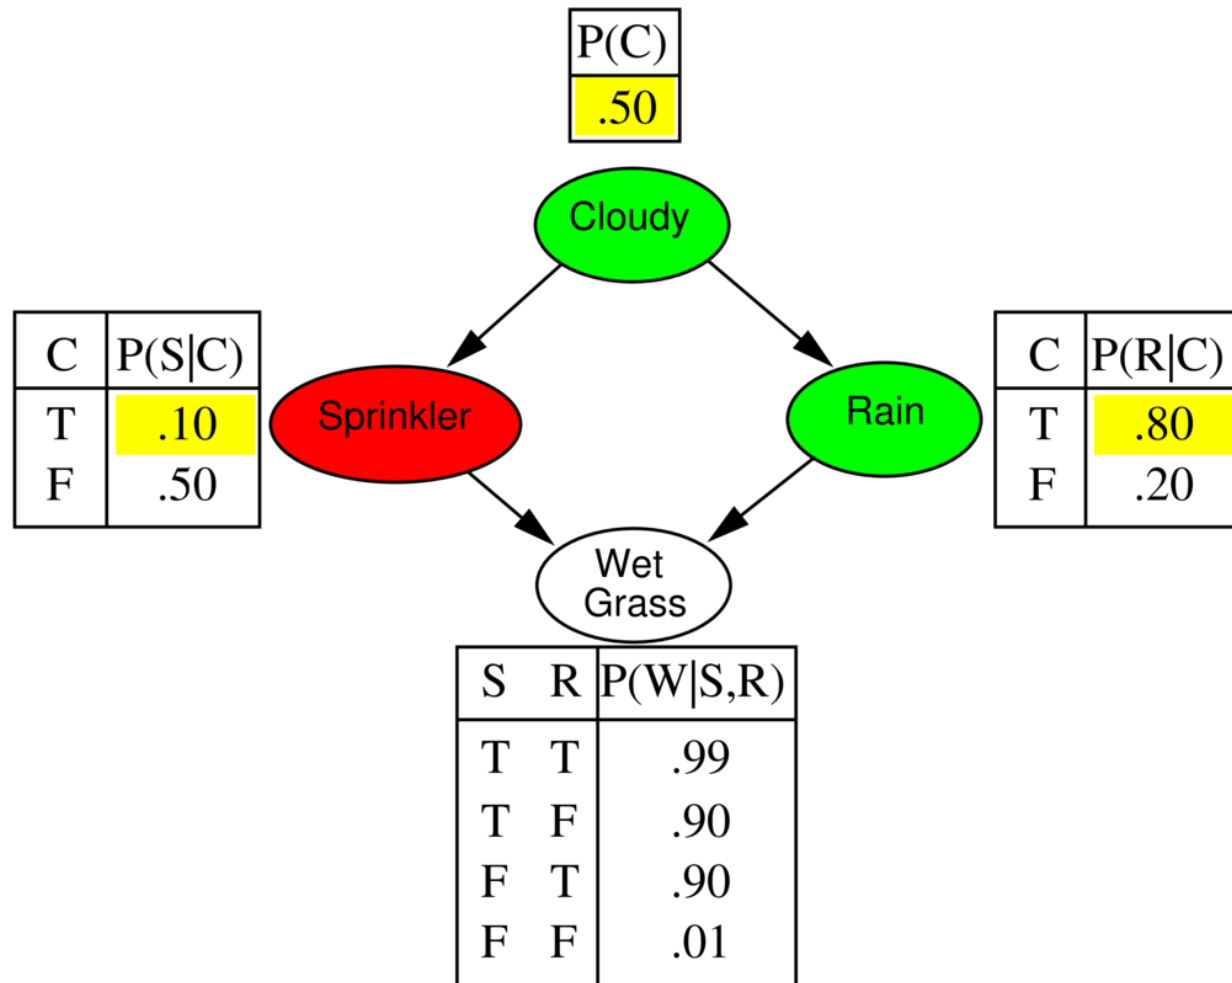


# Prior Sampling



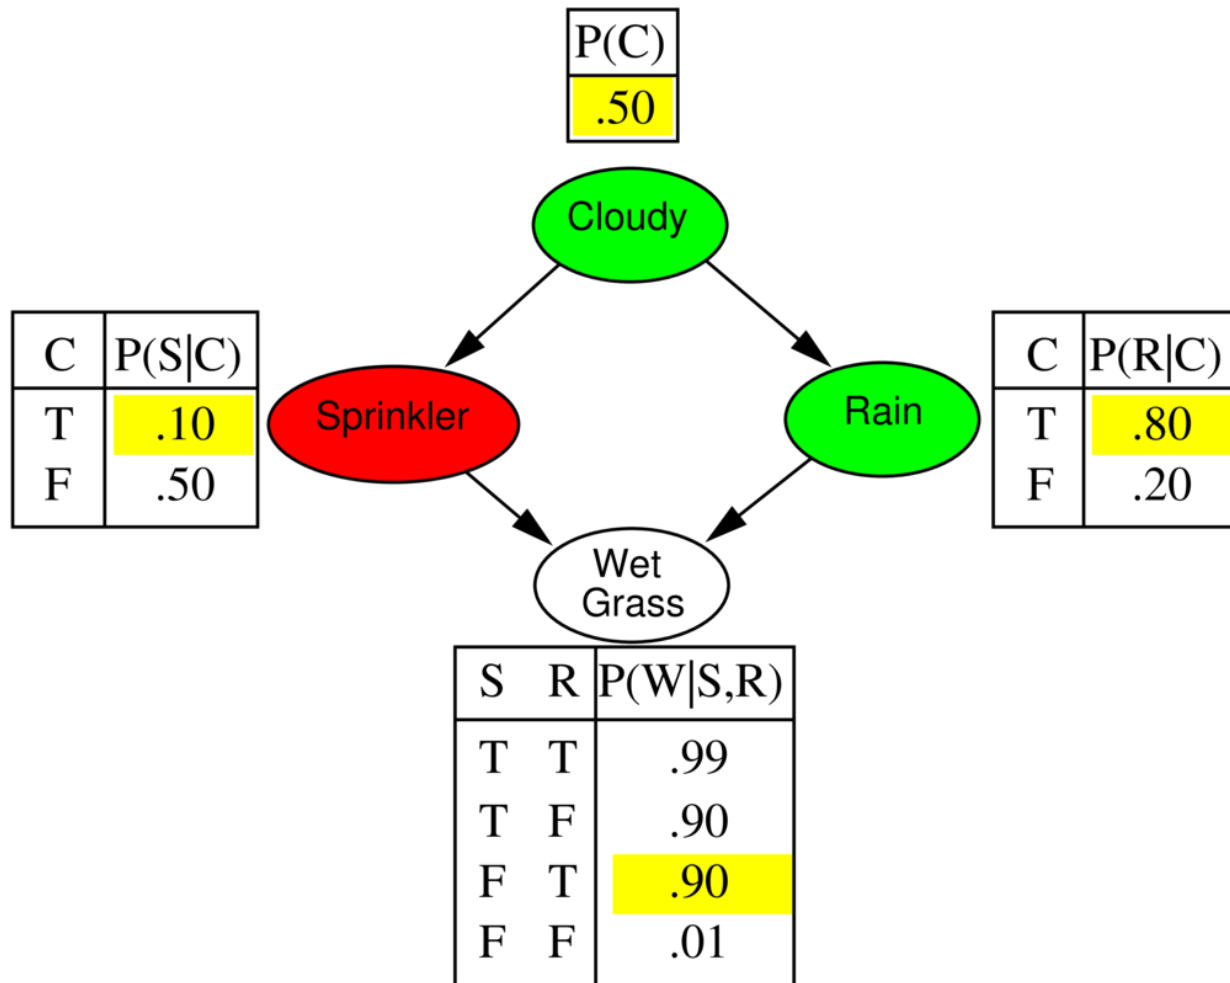
Sprinkler = false

# Prior Sampling

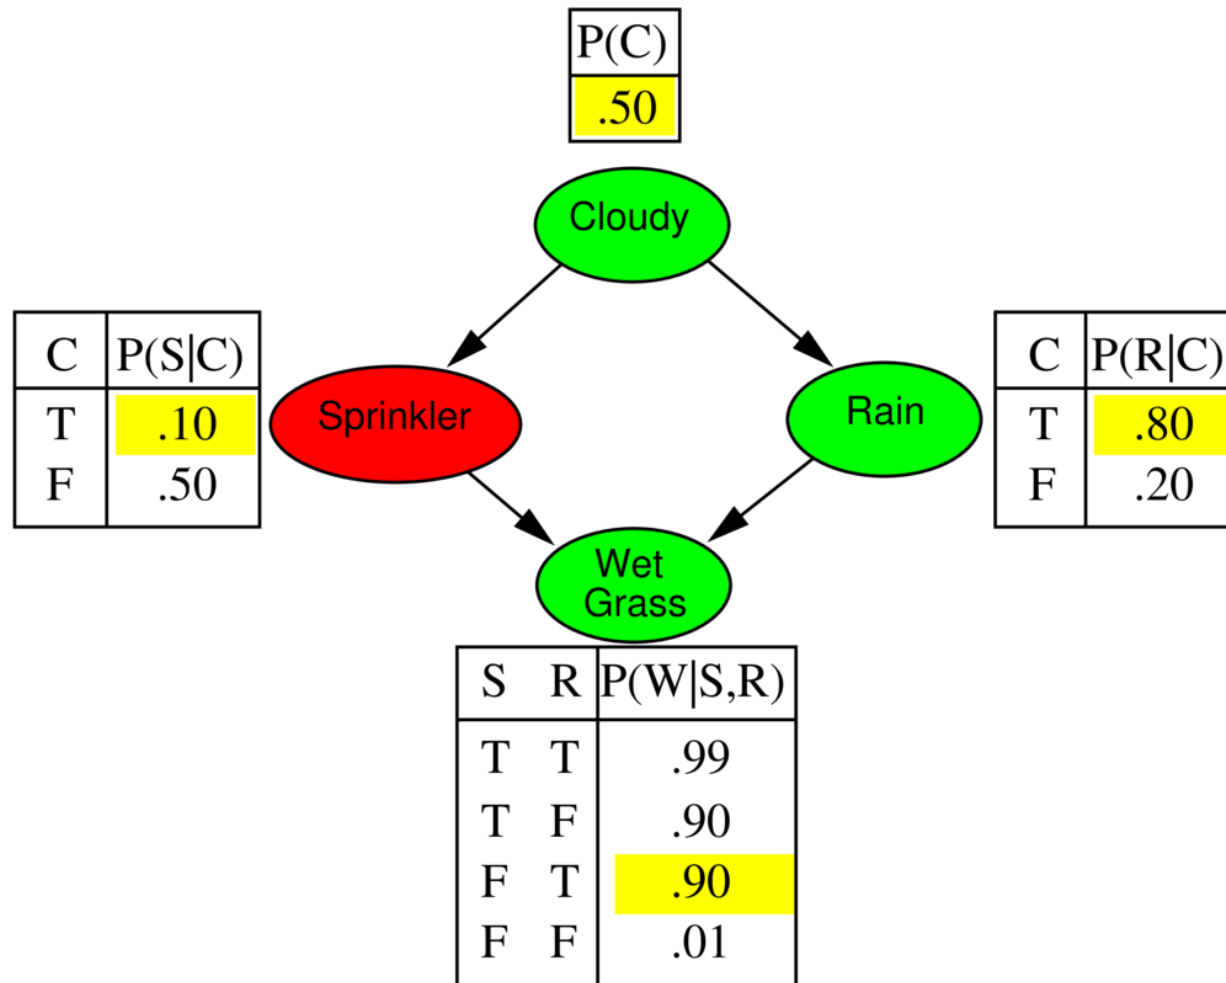


Rain = true

# Prior Sampling



# Prior Sampling



WetGrass = true



# Prior Sampling

- So, we get [*Cloudy = true, Sprinkler = false, Rain = true, WetGrass = true*]
- We write this as [*true, false, true, true*], assuming a specific order of random variables (hierarchical).
- If we repeat the process many times ( $N$ ), we can count the number of times [*true, false, true, true*] is the result ( $m$ ).
- The proportion  $\frac{m}{N}$  gives us  $P(c, \neg s, r, w)$
- The more runs, the more accurate the probability.

# Prior Sampling: Question for you

- Similarly, the proportion  $[false, true, true, true]$  gives us  $P(\neg c, s, r, w)$ . Thus, given:
  - $[true, false, true, true]$
  - $[false, false, true, true]$
  - $[true, true, true, true]$
  - $[true, false, false, true]$
  - $[true, false, true, true]$
  - $[true, false, true, false]$
- $P(c, \neg s, r, w) \approx ?$
- $P(\neg c, s, r, w) \approx ?$

# Prior Sampling: Algorithm

**function** **PRIOR-SAMPLE**(*bn*) **returns** an event sampled from *bn*

**inputs:** *bn*, a belief network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$

**x**  $\leftarrow$  an event with  $n$  elements

**for**  $i = 1$  **to**  $n$  **do**

$x_i \leftarrow$  a random sample from  $\mathbf{P}(X_i \mid \text{parents}(X_i))$   
given the values of  $\text{Parents}(X_i)$  in **x**

**return** **x**

# Rejection Sampling: Algorithm

**function** REJECTION-SAMPLING( $X, \mathbf{e}, bn, N$ ) **returns** an estimate of  $P(X|\mathbf{e})$

**local variables:**  $\mathbf{N}$ , a vector of counts over  $X$ , initially zero

**for**  $j = 1$  to  $N$  **do**

$\mathbf{x} \leftarrow \text{PRIOR-SAMPLE}(bn)$

**if**  $\mathbf{x}$  is consistent with  $\mathbf{e}$  **then**

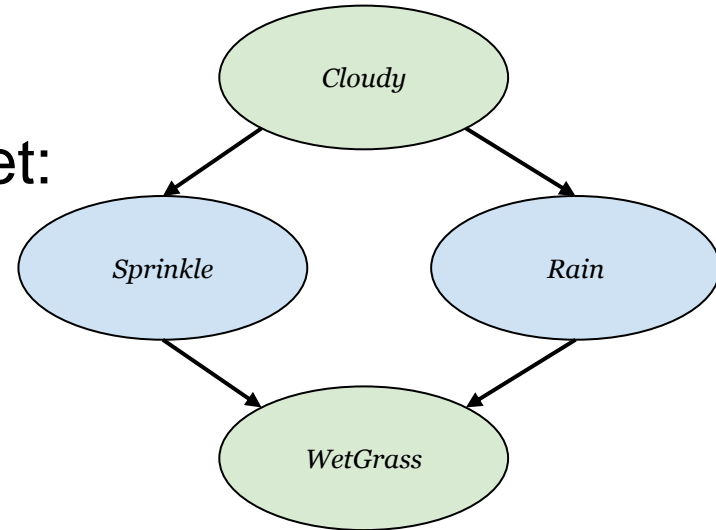
$\mathbf{N}[x] \leftarrow \mathbf{N}[x] + 1$  where  $x$  is the value of  $X$  in  $\mathbf{x}$

**return** NORMALIZE( $\mathbf{N}[X]$ )

Just calls prior sampling, but only counts cases that are consistent with the evidence.

# Rejection Sampling: Example

Consider the Sprinkler Bayes net:



Estimate  $P(\text{Rain} \mid \text{Sprinkler} = \text{true})$  from  $N=100$

- 73 have  $\text{Sprinkler} = \text{false}$ , therefore 73 are **rejected**
- Out of the remaining samples (27) with  $\text{Sprinkler} = \text{true}$ , **8** have  $\text{Rain} = \text{true}$  and **19** have  $\text{Rain} = \text{false}$
- The approximate distribution is  
 $\hat{P}(\text{Rain} \mid \text{Sprinkler} = \text{true}) = \alpha < 8, 19 > = < 0.296, 0.704 >$

# Rejection Sampling: Summary

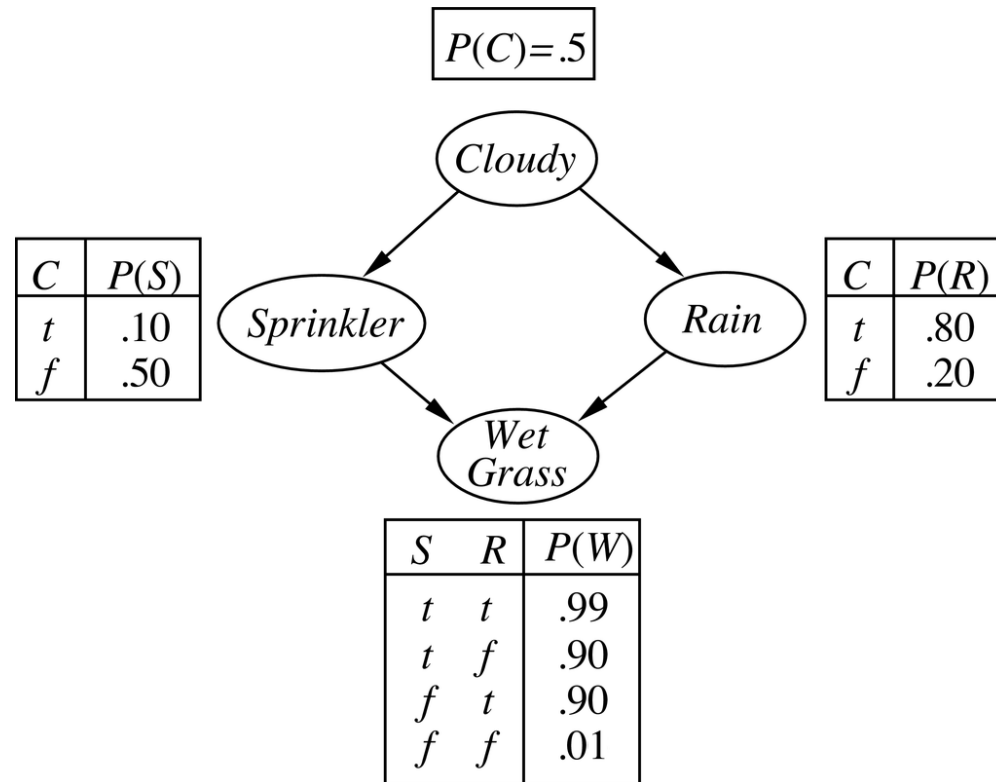
- Rejection sampling uses (conditional) probabilities to sample events, reject those that do not match the evidence, and calculate posterior probabilities.
- For unlikely events, may have to wait a long time to get enough matching samples.
- That can be inefficient.
- If that's the case, use alternative algorithms.

# Today

- Inference by stochastic simulation
- Approximate inference in Bayesian nets (BNs)
  - Rejection Sampling
  - **Likelihood Weighting**
  - Gibbs Sampling
- Data discretisation for learning discrete BNs
- Questions and answers related to assignment

# Likelihood Weighting: Example

Consider we have the Sprinkler network:



Say we want to establish

$$P(Rain | Cloudy = true, WetGrass = true)$$

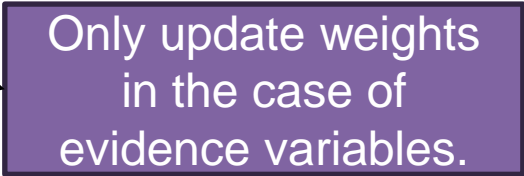


# Likelihood Weighting: Example

- We pick a variable ordering as before, say: *Cloudy, Sprinkler, Rain, WetGrass*.
- Set the weight  $w = 1$  and we start.
- Deal with each variable in order.
- *Cloudy* is **true**, so:
  - $w = w * P(\textit{Cloudy} = \textit{true}) = 1 * 0.5 = 0.5$
- Event:
  - [*Cloudy* = *true*, *Sprinkler* =?, *Rain* =?, *WetGrass* =?]

# Likelihood Weighting: Example

- *Sprinkler* is not an evidence variable, so we do not know whether it is true or false.
- Sample a value just as we did for prior sampling:
  - $P(\textit{Sprinkler} | \textit{Cloudy} = \textit{true}) = \langle 0.1, 0.9 \rangle$
- Let's assume this returns **false**.
- $w$  remains the same.
- Event:
  - [*Cloudy* = *true*, *Sprinkler* = *false*, *Rain* = ?, *WetGrass* = ?]



Only update weights  
in the case of  
evidence variables.

# Likelihood Weighting: Example

- *Rain* is not an evidence variable, so we also do not know whether it is true or false.
- Sample a value just as we did for prior sampling:
  - $P(Rain|Cloudy = true) = \langle 0.8, 0.2 \rangle$
- Let's assume this returns **true**.
- $w$  remains the same.
- Event:
  - [*Cloudy* = *true*, *Sprinkler* = *false*, *Rain* = *true*, *WetGrass* = ?]

# Likelihood Weighting: Example

- *WetGrass* is an evidence variable, with value true, so we set:
  - $w = w * P(WetGrass = true | Sprinkler = false, Rain = true) = 0.5 * 0.90 = 0.45$
- Event:
  - [*Cloudy* = true, *Sprinkler* = false, *Rain* = true, *WetGrass* = true]
- So we end up with event [*true*, *false*, *true*, *true*] and weight  $w = 0.45$ .

# Likelihood Weighting: Example

10 sampled events:

<i>Cloudy</i>	<i>Sprinkler</i>	<i>Rain</i>	<i>GrassWet</i>	<i>w</i>
<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	0.45
<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>	0.475
<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	0.05
<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	0.45
<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	0.45
<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	0.45
<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>	0.475
<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	0.45
<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	0.05
<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	0.45

To find  $P(\text{Rain} = \text{true} | \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$  we add up the green-related weights and divide by the sum of all the weights.

# Likelihood Weighting: Full Example

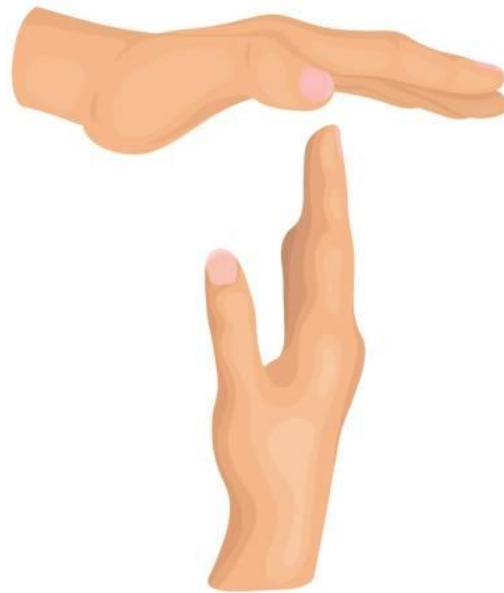
Query  $P(R \mid c, wg)$ , assuming order [*Cloudy*, *Sprinkler*, *Rain*, *WetGrass*]

- $w \leftarrow 1$
- *Cloudy* is evidence, therefore  $w = w \times P(c) = 1 \times 0.5 = 0.5$
- *Sprinkler* is not evidence, therefore sample from  $P(S \mid c)$  — assume *false*
- *Rain* is not evidence, therefore sample from  $P(R \mid c)$  — assume *true*
- *WetGrass* is evidence, thus  $w = w \times P(wg \mid \neg s, r) = 0.5 \times 0.9 = 0.45$
- Return event [*true*, *false*, *true*, *true*] with weight  $w = 0.45$

Repeating that 10 times:

$$\begin{aligned}
 P(r|c, g) &= \frac{\sum_r w}{\sum w} \\
 &= \frac{0.45 + 0.475 + 0.45 + 0.45 + 0.45 + 0.475 + 0.45 + 0.45}{0.45 + 0.475 + 0.05 + 0.45 + 0.45 + 0.45 + 0.475 + 0.45 + 0.05 + 0.45} \\
 P(\neg r|c, g) &= \frac{\sum_{\neg r} w}{\sum w} \\
 &= \frac{0.05 + 0.05}{0.45 + 0.475 + 0.05 + 0.45 + 0.45 + 0.45 + 0.475 + 0.45 + 0.05 + 0.45} \\
 &= 0.0267
 \end{aligned}$$

# Break



# Likelihood Weighting: Algorithm

**function** LIKELIHOOD-WEIGHTING( $X, \mathbf{e}, bn, N$ ) **returns** an estimate of  $P(X|\mathbf{e})$

**local variables:**  $\mathbf{W}$ , a vector of weighted counts over  $X$ , initially zero

**for**  $j = 1$  to  $N$  **do**

$\mathbf{x}, w \leftarrow \text{WEIGHTED-SAMPLE}(bn)$

$\mathbf{W}[x] \leftarrow \mathbf{W}[x] + w$  where  $x$  is the value of  $X$  in  $\mathbf{x}$

**return** NORMALIZE( $\mathbf{W}[X]$ )

Adds in a new weight for a value of  $X$ .



# Likelihood Weighting: Algorithm

**function** **WEIGHTED-SAMPLE**( $bn, \mathbf{e}$ ) **returns** an event and a weight

$\mathbf{x} \leftarrow$  an event with  $n$  elements;  $w \leftarrow 1$

**for**  $i = 1$  **to**  $n$  **do**

**if**  $X_i$  has a value  $x_i$  in  $\mathbf{e}$

**then**  $w \leftarrow w \times P(X_i = x_i \mid \text{parents}(X_i))$

**else**  $x_i \leftarrow$  a random sample from  $\mathbf{P}(X_i \mid \text{parents}(X_i))$

**return**  $\mathbf{x}, w$

Generates a new sample.

# Today

- Inference by stochastic simulation
- Approximate inference in Bayesian nets (BNs)
  - Rejection Sampling
  - Likelihood Weighting
  - **Gibbs Sampling**
- Data discretisation for learning discrete BNs
- Questions and answers related to assignment

# Gibbs Sampling

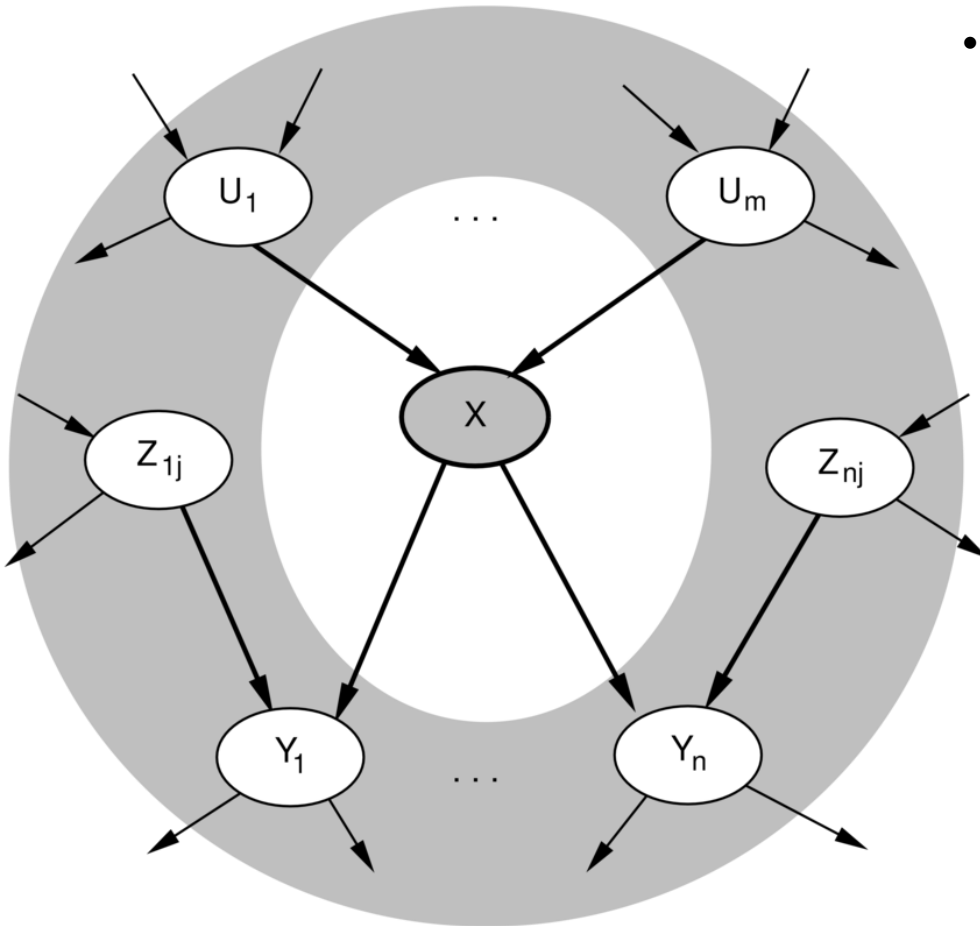
- A rather different approach to sampling.
- Part of the Markov Chain Monte Carlo (MCMC) family of algorithms.
- **Idea:** Don't generate each sample from scratch.
- Generate samples by making a random change to the previous sample.

# Gibbs Sampling

- Gibbs sampling starts with an **arbitrary state**.
- So, pick state with **evidence variables fixed** at observed values (if *Cloudy=true*, we pick that).
- Generate **next state** by randomly sampling from a non-evidence variable – but conditional on the current values of the **Markov blanket**.
- “The algorithm therefore wanders randomly around the state space . . . flipping one variable at a time, but keeping evidence variables fixed”.

# Markov Blanket

- The Markov blanket of  $X$  is:
  - Its parents ( $U_i$ );
  - Its children ( $Y_i$ ); and
  - Its children's (other) parents ( $Z_i$ ).



# Gibbs Sampling: Example

- Start with a query:
  - $P(Rain | Sprinkler = true, WetGrass = true)$
- Our initial state takes values of evidence variables:
  - $[Cloudy = ?, Sprinkler = true, Rain = ?, WetGrass = true]$and the other two are initialised randomly, e.g.:
  - $[Cloudy = true, Sprinkler = true, Rain = false, WetGrass = true]$
- Now we can start the main sampling process.

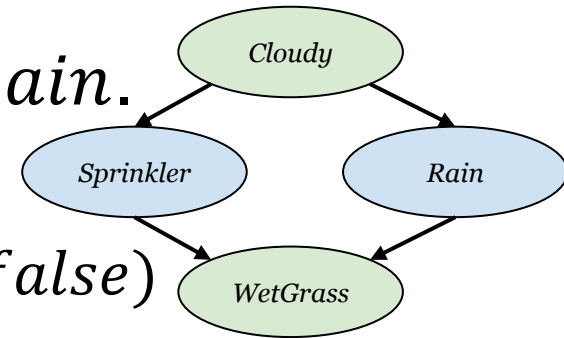
# Gibbs Sampling: Example

- First we sample *Cloudy* given the current state of its Markov blanket.

- Markov blanket is *Sprinkler* and *Rain*.

- So, sample from:

- $P(\textit{Cloudy} | \textit{Sprinkler} = \textit{true}, \textit{Rain} = \textit{false})$



- Suppose we get *Cloudy* = *false*, then the new state is:
  - [*Cloudy* = *false*, *Sprinkler* = *true*, *Rain* = *false*, *WetGrass* = *true*]

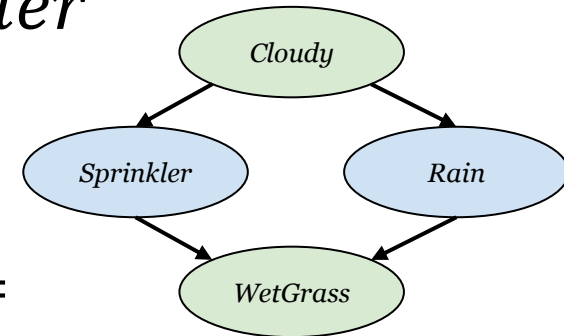
# Gibbs Sampling: Example

- Next we sample *Rain* given its Markov blanket.

- Markov blanket is *Cloudy*, *Sprinkler* and *WetGrass*

- So, sample from:

- $P(\text{Rain} | \text{Cloudy} = \text{false}, \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{false})$



- Suppose we get *Rain* = *true*, then new state is:

- [*Cloudy* = *false*, *Sprinkler* = *true*, *Rain* = *true*, *WetGrass* = *true*]

- Then repeat the sampling.



# Gibbs Sampling: Example

- Each state visited during this process contributes to our estimate for:
  - $P(\textit{Cloudy} | \textit{Sprinkler} = \textit{true}, \textit{WetGrass} = \textit{true})$
- Say the process visits 80 states:
  - In 20,  $\textit{Cloudy} = \textit{true}$
  - In 60,  $\textit{Cloudy} = \textit{false}$
- Then
  - $P(\textit{Cloudy} | \textit{Sprinkler} = \textit{true}, \textit{WetGrass} = \textit{true})$
  - $= \textit{Normalise}(< 20, 60 >)$
  - $= < 0.25, 0.75 >$

# Gibbs Sampling: Algorithm

**function** GIBBS-ASK( $X, \mathbf{e}, bn, N$ ) **returns** an estimate of  $P(X|\mathbf{e})$

**local variables:**  $\mathbf{N}[X]$ , a vector of counts over  $X$ , initially zero

$\mathbf{Z}$ , the nonevidence variables in  $bn$

$\mathbf{x}$ , the current state of the network, initially copied from  $\mathbf{e}$

initialize  $\mathbf{x}$  with random values for the variables in  $\mathbf{Z}$

**for**  $j = 1$  to  $N$  **do**

**for each**  $Z_i$  in  $\mathbf{Z}$  **do**

sample the value of  $Z_i$  in  $\mathbf{x}$  from  $\mathbf{P}(Z_i | mb(Z_i))$   
given the values of  $MB(Z_i)$  in  $\mathbf{x}$

$\mathbf{N}[x] \leftarrow \mathbf{N}[x] + 1$  where  $x$  is the value of  $X$  in  $\mathbf{x}$

**return** NORMALIZE( $\mathbf{N}[X]$ )

# Gibbs Sampling: Markov Blanket

- All of this begs the question: “*How do we sample a variable given the state of its Markov blanket?*”

- For a value  $x$  of a variable  $X$ :

$$P(X|mb(X)) = \alpha P(X|parents(X)) \prod_{Y \in Children(X)} P(Y|parents(Y))$$

where  $mb(X)$  is the Markov blanket of  $X$ .

- Given  $P(X|mb(X))$ , we can sample from it just as we have before.

# Gibbs Sampling: Markov Blanket

- For the sprinkler example, consider we want to compute:
  - $P(\textit{Cloudy} | \textit{Sprinkler} = \textit{true}, \textit{Rain} = \textit{false})$
- Then we need:

$$P(X | mb(X)) = \alpha P(X | \textit{parents}(X)) \prod_{Y \in \textit{Children}(X)} P(Y | \textit{parents}(Y))$$

$$P(C | mb(C)) = \alpha P(C | \textit{parents}(C)) \prod_{Y \in \textit{Children}(C)} P(Y | \textit{parents}(Y))$$

# Gibbs Sampling: Markov Blanket

- Which gives us:

$$\begin{aligned} P(C|mb(C)) &= \alpha P(C|parents(C)) \prod_{Y \in Children(C)} P(Y|parents(Y)) \\ &= \alpha P(C)P(s|C)P(\neg r|C) \\ &= \alpha < P(c), P(\neg c) > < P(s|c), P(s|\neg c) > < P(\neg r|c), P(\neg r|\neg c) > \\ &= \alpha < 0.5, 0.5 > < 0.1, 0.5 > < 0.2, 0.8 > \\ &= \alpha < 0.5 * 0.1 * 0.2, 0.5 * 0.5 * 0.8 > \\ &= \alpha < 0.01, 0.2 > \\ &= < 0.048, 0.952 > \end{aligned}$$

- And then we sample as normal.

# Ideas for your Assignment (optional)

1. Implement the Likelihood Weighting algorithm (*by extending BayesNetInference.py*)
  2. Implement the Gibbs Sampling algorithm in the code of today's workshop (*by extending BayesNetInference.py and related files*)
- ...

But give priority to solving the assignment task first using the code provided during the workshops. Once you have done that, think about implementing your own algs.

# Today

- Inference by stochastic simulation
- Approximate inference in Bayesian nets (BNs)
  - Rejection Sampling
  - Likelihood Weighting
  - Gibbs Sampling
- **Data discretisation for learning discrete BNs**
- Questions and answers related to assignment

# How to train discrete Bayes Nets with Continuous Random Variables?

1. Manual discretisation of continuous data—to be able to use the methods discussed so far.
  - Age (years/days/minutes): young, middle aged, old
  - Temperature: very cold, cold, mild, hot, very hot
2. Split continuous data into uniform intervals (e.g., 1, 2, 3,... intervals between min & max ).
3. Data discretisation using *bnlearn*\*.

\* Chen, Y., et al. "[Learning Discrete Bayesian Networks from Continuous Data](#)", JAIR, 2017



# API for Data Discretisation

[bnlearn.discretize](#)(data, edges,  
continuous\_columns, max\_iterations, verbose)

Example code snippet (using cardiovascular\_data):

```
data = pd.read_csv(TRAINING_DATA, encoding='latin')
edges = [('target', 'i»age'), ('target', 'gender'), ('target',
'height'), ('target', 'weight'), ('target', 'ap_hi'), ('target',
'ap_lo'), ('target', 'cholesterol'), ('target', 'gluc'),
('target', 'smoke'), ('target', 'alco'), ('target', 'active')]
continuous_columns = ["i»age", "height", "weight", "ap_hi",
"ap_lo"]
df_discrete = bn.discretize(data, edges, continuous_columns,
max_iterations=1, verbose=3)
```

# Data Discretisation: Example Output

```
(env_bnlearn) C:\Lincoln\slides\CMP9794M-2024-25\aa-i-workshop-w4>python bnlearn_DataDiscretisation.py
DATA:
i>age gender height weight ap_hi ap_lo cholesterol gluc smoke alco active target
0 17623 2 169 82 150 100 1 1 0 0 1 1
1 17474 1 156 56 100 60 1 1 0 0 0 0
2 14791 2 165 60 120 80 1 1 0 0 0 0
3 17482 1 154 68 100 70 1 1 0 0 0 0
4 21413 1 157 69 130 80 1 1 0 0 1 0
.. ...
495 21111 1 160 71 140 90 1 2 0 0 1 0
496 19655 1 168 82 120 90 1 1 0 0 1 0
497 19594 2 174 54 120 80 2 2 1 0 1 0
498 19711 1 162 100 120 80 3 1 0 0 1 1
499 21242 2 168 79 130 90 1 1 0 0 1 0

[500 rows x 12 columns]
[bnlearn] >Discretizer for continuous values. Iteration [0].
i>age gender height weight ap_hi ap_lo cholesterol
0 (14213.54, 19994.0] 2 (140.45, 190.0] (79.5, 151.5] (138.5, 563.0] (88.5, 575.0]
1 (14213.54, 19994.0] 1 (140.45, 190.0] (41.74, 79.5] (2.0499999999999999, 138.5] (-11.0, 88.5]
2 (14213.54, 19994.0] 2 (140.45, 190.0] (41.74, 79.5] (2.0499999999999999, 138.5] (-11.0, 88.5]
3 (14213.54, 19994.0] 1 (140.45, 190.0] (41.74, 79.5] (2.0499999999999999, 138.5] (-11.0, 88.5]
4 (19994.0, 23653.0] 1 (140.45, 190.0] (41.74, 79.5] (2.0499999999999999, 138.5] (-11.0, 88.5]
.. ...
495 (19994.0, 23653.0] 1 (140.45, 190.0] (41.74, 79.5] (138.5, 563.0] (88.5, 575.0]
496 (14213.54, 19994.0] 1 (140.45, 190.0] (79.5, 151.5] (2.0499999999999999, 138.5] (88.5, 575.0]
497 (14213.54, 19994.0] 2 (140.45, 190.0] (41.74, 79.5] (2.0499999999999999, 138.5] (-11.0, 88.5]
498 (14213.54, 19994.0] 1 (140.45, 190.0] (79.5, 151.5] (2.0499999999999999, 138.5] (-11.0, 88.5]
499 (19994.0, 23653.0] 2 (140.45, 190.0] (41.74, 79.5] (2.0499999999999999, 138.5] (88.5, 575.0]
```

Those outputs can be used by `CPT_Generator.py`,  
`BayesNetInference.py`, and `ModelEvaluator.py`

# Today

- Inference by stochastic simulation
- Approximate inference in Bayesian nets (BNs)
  - Rejection Sampling
  - Likelihood Weighting
  - Gibbs Sampling
- Data discretisation for learning discrete BNs
- **Questions and answers related to assignment**

# Assessment Item 1: Assignment

https://blackboard.lincoln.ac.uk/ultra/courses/\_200412\_1/outline

Module Content | Calendar | Announcements | Discussions | Grade Centre | Messages | Analytics | Groups

**Assessment Item 1**  
Visible to students  
This folder contains all of the documentation and upload areas for the first assessment in this module.

**Assessment Item 1 Documentation**  
Visible to students  
Please find attached the briefing document and CRG for assessment item 1 (assignment).

**CMP9794M Advanced Artificial Intelligence Assessment Item 1 Brief 24-25.pdf**  
Visible to students

**CMP9794M Advanced Artificial Intelligence Assessment Item 1 CRG 24-25.pdf**  
Visible to students

**datasets-assessment-item1.zip**  
Visible to students

read

use

# First 4 Weeks of this Semester

We have asked and answered the following when training and evaluating Bayesian Networks (BNs):

1. **Where do the probabilities come from?**  
*A: Usually from training data (MLE, not the only one)*
2. **Where does the structure come from?**  
*A: Prior knowledge or structure learning (preferred)*
3. **How to do probabilistic inference?**  
*A: Inference by enumeration, variable elimination, rejection sampling, likelihood weighting, Gibbs sampling*
4. **What code to use for training/evaluating BNs?**  
*A: Module's code and bnlearn code – you've got both.*

# Workshop Materials in Github

<https://github.com/LCAS/CMP9794M>

## Usage:

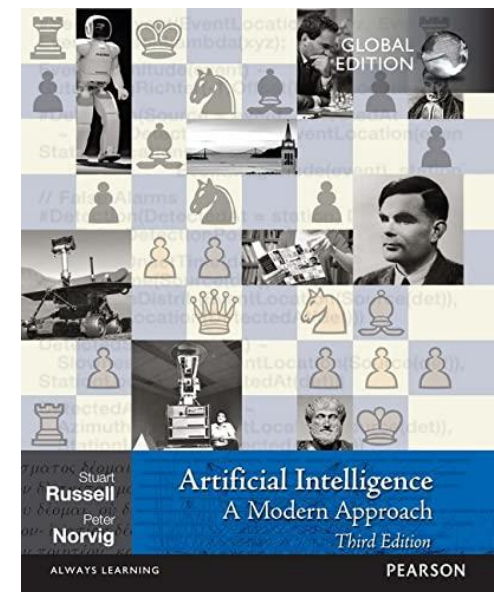
1. Open a terminal and change directory to the path where you want to download the repository.
2. Terminal prompt> git clone <https://github.com/LCAS/CMP9794M>
3. Terminal prompt> cd CMP9794M
4. Terminal prompt> git pull
5. Open VS Code and you should see the workshop folders, one per week.
6. Run the code as done in the workshop sessions.

# Today

- Approximate inference should be used if exact inference becomes unfeasible/impractical.
- We looked at algorithms for approximate inference in Bayes nets—based on sampling.
  - Rejection sampling
  - Likelihood weighting
  - Gibbs sampling

## Readings:

Russell & Norvig 2016. Chapter 14.5



# **This and Next Week**

## **Workshop (today):**

Code for approximate probabilistic inference via  
Rejection Sampling, Gibbs Sampling, etc.

## **Lecture (next week):**

Gaussian Bayesian Networks

Reading 1: Koller & Friedman 2009. [Section 7.2](#)

Reading 2: Bishop. C. 2006. [Section 8.1.4](#)

Any other questions?